

Higher-Level Synchronization Mechanisms

AND-Synchronization
Sequencer and Event Count
Monitor

AND Synchronization

- As we have seen with the **Cigarette Smoker's Problem**, it is often difficult to coordinate access to resources by the means of simple semaphores.
- **Why? What was the problem?**
- Basically, we need to **test for the availability of multiple resources in one step**, i.e., tobacco and matches.
- We can do so with AND-Synch. or parallel semaphores.

We define new semaphore operations, **SP(s)** and **SV(s)**:

SP(s1,t1,d1;s2,t2,d2;.....;sn,tn,dn):

if $s1 \geq t1$ and $s2 \geq t2$ andand $sn \geq tn$ then for all i do $si := si - di$ else place process on queue associated with the first $si < ti$, establish program counter to restart the entire test.

SV(s1,d1;s2,d2;.....;sn,dn)

$si := si + di$ for all i ; and wakeup all processes waiting on any of the si .

Strong Reader Preference

Shared Variables: NR is a semaphore initialized to the total number of readers, R; mx is a semaphore initialized to 1.

Reader

loop

SP(NR,1,1);

SP(mx,1,0);

Perform read;

SV(NR,1);

endloop;

Writer

loop

SP(mx,1,1; NR, R, 0)

Perform write;

SV(mx,1);

endloop;

Writer Preference

Shared Variables: **NW** is a semaphore initialized to W, the total number of writers; **NR** is a semaphore initialized to R, the total number of readers; **mx** is a semaphore initialized to 1.

Writer

loop

SP(NW,1,1);

SP(mx,1,1;NR,R,0);

Perform write;

SV(mx,1; NW,1);

endloop;

Reader

loop

SP(NR,1,1; NW,W,0);

Perform read;

SV(NR,1);

endloop;

Sequencer and Eventcounts

- An interesting mechanism for coordinating processes is the **Sequencer & Eventcount**.
- Consider a bakery or a bank with multiple sales persons/clerks.
- Multiple customers can enter the facility at the same time and must be organized such that they are served in order as soon as service is available.
- This problem has been solved by installing a ticket machine that issues tickets in numerical order.
- As soon as any of the servers become available, a counter that displays a number for the **Next Customer to be Served** is incremented.
- The same mechanism can be used to coordinate and synchronize processes.
- In OS, the stack or of tags by which customers are ordered is represented by an **Eventcount E**.
- The machine that issues tickets or tags to the customers is represented as a **Sequencer S**.

...Sequencer & Eventcount

- Only a single operation is defined on the sequencer S:

- `ticket(S)` :: issues a non-negative, increasing, and contiguous sequence of integers.

- The `ticket(S)` operation corresponds to a newly arriving customer taking a unique numbered tag.

- The eventcount E has 3 associated functions:

- `await(E,v);`
- `advance(E);`
- `read(E);`

A sequence

`v := ticket(s);`
`await(E,v)`

causes the process (customer) to wait until E reaches v.

`await(E,v)` suspends the calling process if $E < v$; otherwise it allows the process to proceed.

We may of course combine the two calls above to

`await(E, ticket(s));`

Sequencer & Eventcount

Formally, $\text{await}(E, v)$ is defined as:

$\text{await}(E, v)$:

if $E < v$ then place the calling process in the queue that is associated with E and invoke the scheduler

$\text{advance}(E)$ corresponds to an initiation of service: the eventcount value E is incremented and the next process/customer is admitted for service.

$\text{advance}(E)$:

$E := E + 1$;

Wakeup the process(es) waiting for E 's value to reach the current value just obtained;

$\text{read}(E)$ provides a means for inspecting the current value of E . This may be useful as a process may want to check how long it may have to wait.

Example: Solving the CS problem

E : eventcount /* initialized to 0 */

S : sequencer /* initialized to 0 */

$\text{await}(E, \text{ticket}(S))$;

enter CS;

$\text{advance}(E)$;

Example: Producer / Consumer

Shared Variables

```
var Pticket, Cticket: sequencer;  
    In, Out: eventcount;  
    buffer: array[0..N-1] of item;
```


!!! Remember, Pticket, Cticket, In, and Out are initialized to 0 !!!

Producer i:

```
var t: integer;  
loop  
    Create new item;  
    t := ticket(Pticket)  
    await(In, t); /* one at a time */  
    await(Out, t-N+1);  
    buffer[t mod N] = item;  
    advance(In);  
endloop;
```

Consumer k:

```
var u: integer;  
  
loop  
    u := ticket(Cticket);  
    await(Out, u); /* one at a time */  
    await(In, u+1);  
    item := buffer[u mod N];  
    advance(Out);  
    consume item ;  
endloop
```



Monitor - an ADT

- One of the disadvantages of any of the synchronization mechanisms seen thus far is that they are very **low-level and hence prone to programming errors**.
- A **MONITOR** is a structure that contains data and functions.
- In other words, a monitor can be viewed as an **Abstract Data Type**, which facilitates the synchronization and coordination of access to objects.
- The advantages of the object oriented paradigm apply.
- The following are the characteristics of a monitor:
 - **Mutual exclusive access** → only one process can be active inside the monitor.
 - Access to any of the encapsulated functions is only possible via the **monitor procedure**.
 - Use of **condition variables (CV)** → not really variables (no memory is associated with it).
 - Monitor semantics is based on **Dijkstra's Guarded Command**, e.g.: **[condition] → action**

..condition variables (CV)

- A condition variable is a name that is chosen by the programmer to represent a specific state or condition. For example:

- condition: CS_Open

may be used to represent the fact that the critical section is currently not used.

- Associated with each CV is a queue which can hold processes that are waiting for the condition represented by CV to become "true".


- The following functions are defined for CVs:

- **CV.wait** → causes the executing process to be suspended on the queue associated with CV.
- **CV.signal** → wakes up a process that waiting on CV, if one exists, otherwise no-op.
- **CV.queue** → true if queue is not empty, false if no process is waiting in the queue.

- We need to revisit the **semantics** after discussing the structure of a monitor.

Queues and more queues

- A monitor maintains many different queues:
 - Entry Queue → sequence processes to ensure mutual exclusive access.
 - CV Queues → store suspended processes that are waiting on the corresponding condition.
 - Urgent Queue → (to implement Hoare semantics) stores processes that have signaled others and are about to exit the monitor.



Monitor Examples...

A single resource allocator with Monitors

```
monitor: single_resource
{
  boolean: busy = false; /* initialization */
  condition: non_busy;

  acquire(){
    if busy
      non_busy.wait;
    busy = true;
  }

  release(){
    busy = false;
    non_busy.signal;
  }
}
```

An **extension** to CV.wait is **CV.wait(p)** where p is a priority number (the smaller p , the higher the priority)

Example: Shortest Job First (SJF)

```
monitor: SJN {
  condition x;
  boolean busy = false;

  acquire(int: time){
    if busy x.wait(time);
    busy = true;
  }

  release(){
    busy = false;
    x.signal;
  }
}
```

Different Semantics

- Recall, that only one process can be active inside the monitor at any given moment!
- We need to consider what happens if one process executes a CV.signal, thereby freeing a suspended process.
- This leads to two different CV.signal semantics:
 - Hoare Semantics
 - Mesa Semantics
- In the Hoare semantics
 - the process issuing the CV.signal is placed on the urgent queue if a process was waiting on the corresponding CV-queue.
 - it thereby yields the awakened process (only one process must be active).
 - Upon exiting the monitor, a process signals (V(urgent)) if there are processes suspended in the urgent queue.
 - If the urgent queue is empty, the exiting process will release access to the monitor.

...more semantics...

■ In the **Mesa** Semantics:

- the process issuing a **CV.signal** causes a waiting process to be placed on a ready queue.
- **CV.signal** actually becomes a **notify** operation.
- In this scheme, we cannot guarantee that a particular condition that was signaled to be true, is still true at the time the waiting process will execute.
- Hence, a waiting process must re-evaluate the condition and suspend itself if it is found false.

- A process can easily re-evaluate the condition by using a while-construct instead of a simple if-construct to test the condition.

`while(!B) c.wait;`

instead of the traditional

`if(!B) c.wait;`

The Mesa semantics leads to a more efficient implementation and reduces the number of context switches.

Implementing a Monitor

- The following is an implementation of a Monitor ADT by means of semaphores:

```
P(mutex)
//only one process is active in the monitor
...
body of monitor procedures
...
if next.count > 0 V(next);
else V(mutex);

// where next is a semaphore that
// represents the urgent queue !!
```

CV.wait: →

```
CV.count ++;
if next.count > 0 V(next);
else V(mutex);
P(CV_sem);
CV.count --;
```

CV.signal: →

```
if CV.count > 0 {
    next.count ++;
    V(CV_sem);
    P(next);
    next.count --;
}
```

Strong Reader Pref. with Monitor

The following is a solution to the strong readers priority using semaphores:

```
monitor: readers_writers;  
int read_count;  
boolean: busy = false  
condition: oktoread, oktowrite;  
{
```

```
startread() {  
    if (busy) oktoread.wait;  
    readcount ++;  
    oktoread.signal;  
}
```

```
endread() {  
    readcount --;  
    if (readcount == 0) oktowrite.signal;  
}
```

```
startwrite() {  
    if (readcount > 0 or busy)  
        oktowrite.wait;  
    busy = true;  
}
```

```
endwrite() {  
    busy = false;  
    if (oktoread.queue) oktoread.signal;  
    else oktowrite.signal;  
}
```

```
}
```


... some Monitor exercises

- Try to develop monitor solutions for the following problems:

- Dining Philosophers
- Bounded Buffer
- Cigarette Smokers

- Note: Monitor-based solution usually follow the same style →

- test and wait on *CV* are usually the first statements in a monitor procedure.
- signal or notify are usually the last actions that are performed.