

Deadlock

Basic Definitions

Different Resource Types

Graph Theoretical Treatment

Detection, Prevention, and Avoidance

The Deadlock Problem

Definition: Deadlock is the **permanent** blocking of a set of processes that either compete for system resources or communicate with each other!

Formal treatment of deadlocks goes back over 30 years:

- Dijkstra (1968)
- Holt (1971, 1972)

We distinguish 3 different types of approaches to the deadlock problem:

1. **Detection and Recovery;**
2. **Prevention;**
3. **Avoidance;**

We consider two different types of system resources:

1. **Reusable Resources** → permanent objects with two properties:
 - a) # of units is constant
 - b) each unit is available or allocated to exactly one process;
2. **Consumable Resources** → produced and consumed dynamically:
 - a) # of units will vary over time;
 - b) system can create new units
 - c) process may request, acquire, and consume resource units

Examples

Examples of reusable resources are:

- Devices → printers, tapes, disk, etc.
- Memory
- Service Routines (critical sections!!)

Examples of consumable resources are:

- messages;
- signals;
- interrupts;
- events;
- data structures (dynamic);

Deadlock Example 1:

The diagram illustrates a race condition between two processes, P1 and P2, for two resources: Disk and Tape. The processes are shown as vertical columns of code snippets. P1's code is on the left, and P2's code is on the right. Arrows indicate the flow of execution and resource allocation.

P1:

- ...
- request(Disk);
- ...
- request(Tape);
- ...
- release(Tape);
- release(Disk);
- ...

P2:

- ...
- request(Tape);
- ...
- request(Disk);
- ...
- release(Disk);
- release(Tape);
- ...

Arrows show the following sequence of events:

- P1 requests the Disk.
- P2 requests the Tape.
- P1 requests the Tape.
- P2 requests the Disk.
- P1 releases the Tape.
- P2 releases the Disk.
- P1 releases the Disk.
- P2 releases the Tape.

The arrows cross, indicating that the processes are interleaved in a way that leads to a race condition. Specifically, P1 requests the Tape before P2 releases the Disk, and P2 requests the Disk before P1 releases the Tape.

Let's analyze the problem...

What causes the deadlock condition?

... more examples

Deadlock Example 2:

Available: 100 blocks of memory!

P1:	P2:
...	...
request(40Blocks);	request(30Blocks);
...
request(40Blocks);	request(50Blocks);

Even though neither P1 nor P2 request more than the available amount, deadlock can occur.

Note: their first request can easily be granted!

Deadlock Example 3:

Consider a message exchange between P1 and P2.

P1:	P2:
...	...
receive(P2, M1);	receive(P1, M2));
...	...
send(P2, M1');	send(P1, M2');

Note: Messages are consumable, yet both processes are waiting for a message to be sent by the other process.

Processes and States...

The following actions by P1 and P2 will result in state changes.

	P1:	P2:
0:
	while(1){	while(1){
1:	Req(D);	Req(T);
2:
3:	Req(T);	Req(D);
4:
	Rel(T);	Rel(D);
5:
	Rel(D);	Rel(T);
	}	}


P1:

- 0: Holds none
- 1: Holds none, requests D
- 2: Holds D
- 3: Holds D, requests T
- 4: Holds D and T
- 5: Holds D, T is released

P2:

- 0: Holds none
- 1: Holds none, requests T
- 2: Holds T
- 3: Holds T, requests D
- 4: Holds T and D
- 5: Holds T, D is released

A State Diagram



Note:
There are no
secure states!

Why not ??



A few definitions

- Let Σ be the set of all possible resource allocation states i.e., $\Sigma = \{S, T, U, V\}$
- Let Π be the set of Processes, e.g., $\Pi = \{P_1, P_2\}$
- *A process P_i is blocked in state S* if there does not exist a state T , such that $S \rightarrow^i T$. i.e., no action by processes P_i will result in a state change.
- *Process P_i is deadlocked* in state S if P_i is blocked in state S and for all states T with $S \rightarrow^* T$, P_i is blocked in T
- If P_i is deadlocked in state S , then S is a *deadlock state*.
- If all processes P_i are deadlocked in state S , Then S is a *total deadlock state*.
- *State S is secure* if S is not a deadlock state and, any state T reachable from S ($S \rightarrow^* T$) is not a deadlock state.
- **NOTE:** it is possible for a process P_i to be deadlocked while other processes are running.

Necessary Deadlock Conditions

- The following are necessary conditions that must exist for deadlock to occur:

1. **Mutual Exclusion** → Processes hold resources exclusively, hence making them unavailable to other processes.
2. **Nonpreemption** → Resources cannot be taken away from a process that is holding them. Only the process can release the resources they hold.

3. **Resource Waiting** → Processes that request unavailable resources (or units) will block until they become available. (Circular Wait)
4. **Partial Allocation** → Processes may hold some resources when they request additional units of the same resource or other resources. (Hold and Wait)

Note: the violation of any one of these conditions will result in a deadlock-free system.

Formal Treatment of Deadlocks

A useful model for the interaction between processes and resources is based on concepts from **Graph Theory**.

We define a **Resource Graph**, $G(V, E)$, where V is the set of vertices and E is the set of edges.

The set of vertices (V) consists of process nodes and resource nodes.


The set of edges (E) consists of

- Request Edges
- Assignment Edges
- Producer Edges


A resource graph, G , is a **directed graph (digraph)** → edges are directed from processes to resources:

- **Request Edges** → from process to resources, signify that a process is requesting that resource (or units thereof).
- **Assignment Edges** → from resource to process, indicating that a particular resource is assigned to a specific process.
- **Producer Edge** → connects consumable resources to processes that produce them.

examples



Request Edge




Assignment Edge



Producer Edge

State Transitions

We can use the resource graph to depict state transition as show in the following example:



$S_0 \xrightarrow{\text{request } 2 \times R \text{ by } P_1} S_1 \xrightarrow{\text{acquire } 2 \times R \text{ by } P_1} S_2 \xrightarrow{\text{release } 1 \times R \text{ by } P_1} S_3$

Requests and Acquisitions

1. **Requests** → Any process p_i that has no outstanding requests (i.e., it is not blocked) in a given state may request units of any number of resources R_j, R_k, \dots
 - In the resource graph, this is reflected by adding edges $(P_i, R_j), (P_i, R_k), \dots$ corresponding to the number of units of each resource requested.
 - For reusable resources, the number of requested must never exceed the number of available resource units.
2. **Acquisition** → A process p_i may acquire resources that have previously been requested, provided that the number of requested units are available.
 - In the resource graph, this is reflected by reversing the request edges (P_i, R_j) to (R_j, P_i) for **reusable resources**.
 - For **consumable resources**, we remove the request edges (P_i, R_j) , and remove the number of resource units from R , indicating that they have been consumed.